

# YUANZHE HU

Homepage: [hust-ai-hyz.github.io](https://hust-ai-hyz.github.io)

Phone: (+1) 619-937-1812 ◊ Email: [yzhu.ml@outlook.com](mailto:yzhu.ml@outlook.com)

## EDUCATION

<b>University of California, San Diego (UCSD)</b> MS. in Computer Science Engineering	<i>Sep. 2024 – Mar. 2026 (expected)</i>
<b>Huazhong University of Science and Technology (HUST)</b> B.S. in Artificial Intelligence (Honored Class, Qiming School) GPA: 3.91/4.0	<i>Sep. 2020 – Jun. 2024</i>

## RESEARCH SUMMARY

**Mathematical Foundations & Optimization for LLMs/SciML:** Leveraged mathematical analysis to investigate training dynamics and generalization. Designed advanced algorithms for efficient LLM compression and training. Previous works include FARMS (ICML 2025)[1] for layer-wise pruning, Model Balancing[2] (EMNLP 2024 Oral) for low-resource fine-tuning, and an ongoing project about analysis of multi-regime dynamics in SciML models targeting ICML 2026.

**Enhancing Memory and Reasoning in LLMs and Agents:** Focused on enhancing long-term history processing and post-training reasoning. Previous works include M+ (ICML 2025) [3] for long-term information retention, K2-Think (Tech Report) [4] for large-scale reasoning, MIRIX (Open-source framework, 3K+ GitHub stars) for multi-agent memory systems, MemoryAgentBench (Under Review) [5] for comprehensive evaluation, and Mem-alpha (Under Review) [6] for RL-based memory management.

## PUBLICATIONS

- [1] **Yuanzhe Hu**, K. Goel, V. Killiakov, and Y. Yang, “Eigenspectrum analysis of neural networks without aspect ratio bias,” in *ICML*, 2025.
- [2] Z. Liu\*, **Yuanzhe Hu\***, T. Pang, Y. Zhou, P. Ren, and Y. Yang, “Model balancing helps low-data training and fine-tuning,” in *EMNLP (Oral Presentation)*, 2024.
- [3] Y. Wang, D. Krotov, **Yuanzhe Hu**, Y. Gao, W. Zhou, J. McAuley, D. Gutfreund, R. Feris, and Z. He, “M+: Extending memorylm with scalable long-term memory,” in *ICML*, 2025.
- [4] LLM 360 Team, Institute of Foundation Models, MBZUAI, “K2-think: A parameter-efficient reasoning system,” in *Tech Report*, 2025.
- [5] **Yuanzhe Hu\***, Y. Wang\*, and J. McAuley, “Evaluating memory in llm agents via incremental multi-turn interactions,” in *Under Review*, 2025.
- [6] Y. Wang, R. Takanobu, Z. Liang, Y. Mao, **Yuanzhe Hu**, J. McAuley, and X. Wu, “Mem- $\alpha$ : Learning memory construction via reinforcement learning,” in *Under Review*, 2025.

\* Equal Contribution

## RESEARCH AND INDUSTRIAL EXPERIENCE

<b>Memory LLM and Agents Benchmarking and Construction</b> [3], [5], [6] <i>CSE Research Course, Supervisors: Yu Wang and Prof. Julian McAuley</i>	Oct 2024 - Now UC San Diego
· Led the development of <b>MemoryAgentBench</b> , a comprehensive benchmark designed to systematically assess the long-term memory of LLM agents via multi-turn interactions, with evaluation criteria based on principles of cognitive science.	

- Constructed and implemented the evaluation framework for M+, and conducted systematic benchmarking against baseline methods.
- Designed the framework for MIRIX's Evaluation on multiple memory agent benchmarks and participated in code maintenance.
- Contributed to Reinforcement Learning(RL) on LLM Agent for long-term memory management, achieving 52% test accuracy on LongMemEval(S\*) with only one-third of the full context window.
- Co-authored three research papers: MemoryAgentBench, RL Memory Agents and the M+ (ICML 2025).

**Empirical Analysis of SFT for Large Reasoning Models [4]** Jun 2025 - Sep 2025  
*Research Intern, Supervisor: Prof. Zhiting Hu* Institute of Foundation Models, MBZUAI/LLM 360

- My role in this project includes data pre-processing, model training, and technical report writing.
- Engineered and built the large-scale supervised fine-tuning (SFT) pipeline for models up to 70B parameters (e.g., LLaMA-3.1-70B, Qwen2.5-32B) on multiple GPU cluster.
- Achieved good results on challenging reasoning benchmarks by leveraging this SFT pipeline, scoring 57.3 on LiveCodeBench (code generation) and 72.1 on AIME 2025 (math reasoning).
- Investigated key issues in SFT, including training loss behavior, data selection and mixing strategies, and the relationship between training dynamics and model performance.

**De-biased LLM Pruning Based on Eigenspectrum and MP-Law [1]** Nov 2024 - Apr 2025  
*Supervisor: Prof. Yaoqing Yang* Dartmouth College

- Theorized and empirically validated that existing eigenspectrum analysis is biased by weight matrix **aspect ratios**, a phenomenon explained by the **Marchenko-Pastur (MP) law**, leading to inaccurate layer diagnostics.
- Developed **FARMS**, a novel subsampling method grounded in **Random Matrix Theory** (RMT), to normalize weight matrices to a fixed aspect ratio, enabling an unbiased, size-invariant evaluation of layer training quality.
- Validated the method's effectiveness across diverse domains (**LLM Pruning, CV, SciML**), consistently **outperforming state-of-the-art** layer-wise optimization methods.
- First-authored a research paper accepted to the **ICML 2025**.

**Layer-wise Optimization for Low-Resource SFT of LLM [2]** Jun 2023 - Jun 2024  
*Supervisors: Dr. Pu Ren and Prof. Yaoqing Yang* Dartmouth College

- Utilized the **spectral analysis of model weights**, based on Heavy-Tailed Self-Regularization (HT-SR) theory, to establish a quantitative link between data scarcity and imbalanced layer-wise training, providing a theoretical foundation for designing superior model diagnosis metrics and optimization methods.
- Developed a dynamic **layer-wise learning rate** scheduling algorithm to rebalance training quality across layers, overcoming limitations of prior optimizers in NLP training scenarios.
- Validated the method's effectiveness through extensive experiments on diverse models and benchmarks, improving LLM test accuracy by **2-10%** in low-data SFT and leading to a **co-first** authored publication at EMNLP 2024 (**Oral Presentation**).

## COMPETITION AWARDS / SERVICE

---

Reviewer for ICLR 2026, ARR 2025 (July / Oct.), and Workshops.  
 Volunteer for ICML 2025

## TECHNICAL SKILLS

---

<b>Programming Languages</b>	Python, C/C++, SQL, Bash, Matlab
<b>Machine Learning Tools</b>	PyTorch, Hugging Face Transformers, LLaMA Factory, Verl