

YUANZHE HU

Homepage: hust-ai-hyz.github.io

Phone: (+1) 619-937-1812 ◊ Email: yzhu.ml@outlook.com

EDUCATION

University of California, San Diego (UCSD)
MS. in Computer Science Engineering

Sep. 2024 – Mar. 2026 (expected)

Huazhong University of Science and Technology (HUST)
B.S. in Artificial Intelligence (Honored Class, Qiming School)
GPA: 3.91/4.0

Sep. 2020 – Jun. 2024

RESEARCH SUMMARY

Mathematical Foundations & Optimization for LLMs/SciML: Leveraged mathematical analysis to investigate training dynamics and generalization. Designed advanced algorithms for efficient LLM compression and training. Previous works include FARMS (ICML 2025)[1] for LLM layer-wise pruning, Model Balancing[2] (EMNLP 2024 Oral) for low-resource fine-tuning, Analysis of multi-regime dynamics in SciML models[3], and using spectrum signals in LLM Clustering[4].

Enhancing Memory and Reasoning in LLMs and Agents: Focused on enhancing long-term history processing and post-training reasoning. My previous works include MemoryAgentBench (ICLR 2026) [5] and MemoryAreana[6] for AI Agent memory comprehensive evaluation, M+ (ICML 2025) [7] for long-term information retention, K2-Think (Tech Report) [8] for large-scale reasoning, MIRIX (Open-source framework, 3K+ GitHub stars) for multi-agent memory systems, and Mem-alpha (Under Review) [9] for RL-based memory management.

PUBLICATIONS AND PREPRINTS

- [1] **Yuanzhe Hu**, K. Goel, V. Killiakov, and Y. Yang, “Eigenspectrum analysis of neural networks without aspect ratio bias,” in *ICML*, 2025.
- [2] Z. Liu*, **Yuanzhe Hu***, T. Pang, Y. Zhou, P. Ren, and Y. Yang, “Model balancing helps low-data training and fine-tuning,” in *EMNLP (Oral Presentation)*, 2024.
- [3] **Yuanzhe Hu***, X. Wang*, Y. Wang*, X. Zhong*, H. Lu, T. Pang, M. W. Mahoney, Y. Yan, P. Ren, and Y. Yang, “Unveiling multi-regime patterns in sciml: Distinct failure modes and regime-specific optimization,” in *Under Review (ICML)*, 2026.
- [4] Z. Zhang*, I. V. Prasad*, **Yuanzhe Hu***, Z. Liu, H. Luo, P. Ren, and Y. Yang, “Spectral signatures of large language models,” in *Under Review (KDD)*, 2026.
- [5] **Yuanzhe Hu***, Y. Wang*, and J. McAuley, “Evaluating memory in llm agents via incremental multi-turn interactions,” in *ICLR*, 2026.
- [6] Z. He*, Y. Wang*, C. Zhi*, **Yuanzhe Hu***, T.-P. Chen*, L. Yin, Z. Chen, T. A. Wu, S. Ouyang, Z. Wang, J. Pei, J. McAuley, Y. Choi, and A. Pentland, “Benchmarking agent memory in inter-dependent multi-session agentic tasks,” in *Under Review (ICML)*, 2026.
- [7] Y. Wang, D. Krotov, **Yuanzhe Hu**, Y. Gao, W. Zhou, J. McAuley, D. Gutfreund, R. Feris, and Z. He, “M+: Extending memoryllm with scalable long-term memory,” in *ICML*, 2025.
- [8] LLM 360 Team, Institute of Foundation Models, MBZUAI, “K2-think: A parameter-efficient reasoning system,” in *Tech Report*, 2025.
- [9] Y. Wang, R. Takanobu, Z. Liang, Y. Mao, **Yuanzhe Hu**, J. McAuley, and X. Wu, “Mem- α : Learning memory construction via reinforcement learning,” in *Under Review (ICML)*, 2025.

* Equal Contribution. Underlined papers are accepted by conferences.

RESEARCH AND INDUSTRIAL EXPERIENCE

Memory LLM and Agents Benchmarking and Construction [5]–[7], [9]

CSE Research Course, Supervisors: Yu Wang and Prof. Julian McAuley

Oct 2024 - Now

UC San Diego

- Led the development of **MemoryAgentBench**, a comprehensive benchmark designed to systematically assess the long-term memory of LLM agents via multi-turn interactions, with evaluation criteria based on principles of cognitive science.
- Designed the framework for MIRIX's Evaluation on multiple memory agent benchmarks and participated in code maintenance.
- Contributed to Reinforcement Learning(RL) on LLM Agent for long-term memory management, achieving 52% test accuracy on LongMemEval(S*) with only one-third of the full context window.
- Co-authored four research papers: MemoryAgentBench(ICLR 2026), MemoryArena and RL Memory Agents and the M+ (ICML 2025).

Empirical Analysis of SFT for Large Reasoning Models [8]

Research Collaborator, Supervisor: Prof. Zhiting Hu

Jun 2025 - Sep 2025

Institute of Foundation Models, MBZUAI

- My role in this project includes data pre-processing, model training, and technical report writing.
- Engineered and built the large-scale supervised fine-tuning (SFT) pipeline for models up to 70B parameters (e.g., LLaMA / Qwen) on multiple H200 GPU cluster.
- Leveraged the SFT pipeline to achieve strong results on reasoning benchmarks, while conducting comprehensive analysis on training dynamics, data selection, and loss behaviors.

De-biased LLM Pruning Based on Eigenspectrum and MP-Law [1]

Research Collaborator, Supervisor: Prof. Yaoqing Yang

Nov 2024 - Apr 2025

Dartmouth College

- Developed **FARMS**, a novel subsampling method grounded in **Random Matrix Theory** (RMT), to normalize weight matrices to a fixed aspect ratio, enabling an unbiased, size-invariant evaluation of layer training quality.
- Validated the method's effectiveness across diverse domains (**LLM Pruning**, **CV**, **SciML**), consistently **outperforming state-of-the-art** layer-wise optimization methods.
- First-authored a research paper accepted to the **ICML** 2025.

Layer-wise Optimization for Low-Resource SFT of LLM [2]

Research Collaborator, Supervisors: Dr. Pu Ren and Prof. Yaoqing Yang

Jun 2023 - Jun 2024

Dartmouth College

- Developed a dynamic **layer-wise learning rate** scheduling algorithm to rebalance training quality across layers, overcoming limitations of prior optimizers in LLM SFT scenarios.
- Validated the method's effectiveness through extensive experiments on diverse models and benchmarks, improving LLM test accuracy by **2-10%** in low-data SFT and leading to a **co-first** authored publication at EMNLP 2024 (**Oral Presentation**).

COMPETITION AWARDS / SERVICE

Reviewer for ICLR 2026, ARR 2025 (July / Oct.), and Workshops.

Volunteer for ICML 2025

TECHNICAL SKILLS

Programming Languages

Python, C/C++, SQL, Bash, Matlab

Machine Learning Tools

PyTorch, Hugging Face Transformers, LLaMA Factory, Verl