

YUANZHE HU

Homepage: hust-ai-hyz.github.io

Phone: (+1) 619-937-1812 ◊ Email: yzhu.ml@outlook.com

EDUCATION

University of California, San Diego (UCSD)

Sep. 2024 – Mar. 2026 (expected)

MS. in Computer Science Engineering

Huazhong University of Science and Technology (HUST)

Sep. 2020 – Jun. 2024

B.S. in Artificial Intelligence (Honored Class, Qiming School)

GPA: 3.91/4.0

RESEARCH SUMMARY

Diagnostics and Optimization for Foundation Models: Designed and improved model training dynamics analysis and hyperparameter optimization by applying mathematical tools from Random Matrix Theory. These methods yield more accurate layer-wise hyperparameters for LLM pruning[1] and Foundation Models fine-tuning[2].

Memory & Reasoning for LLM and Agents: Designed MemoryAgentBench[3], a comprehensive benchmark to evaluate memory agents and LLM across multiple core competencies. Contributed to the design of memory-augmented LLMs (M+)[4], Mem-alpha[5], and the SFT role on K2-Think[6].

Research Advisors: Yaoqing Yang, Julian McAuley, Zhiting Hu.

PUBLICATIONS

- [1] **Yuanzhe Hu**, K. Goel, V. Killiakov, and Y. Yang, “Eigenspectrum analysis of neural networks without aspect ratio bias,” in *ICML*, 2025.
- [2] Z. Liu*, **Yuanzhe Hu***, T. Pang, Y. Zhou, P. Ren, and Y. Yang, “Model balancing helps low-data training and fine-tuning,” in *EMNLP (Oral Presentation)*, 2024.
- [3] **Yuanzhe Hu***, Y. Wang*, and J. McAuley, “Evaluating memory in llm agents via incremental multi-turn interactions,” in *ArXiv*, 2025.
- [4] Y. Wang, D. Krotov, **Yuanzhe Hu**, Y. Gao, W. Zhou, J. McAuley, D. Gutfreund, R. Feris, and Z. He, “M+: Extending memoryllm with scalable long-term memory,” in *ICML*, 2025.
- [5] Y. Wang, R. Takanobu, Z. Liang, Y. Mao, **Yuanzhe Hu**, J. McAuley, and X. Wu, “Mem- α : Learning memory construction via reinforcement learning,” in *ArXiv*, 2025.
- [6] LLM 360 Team, Institute of Foundation Models, MBZUAI, “K2-think: A parameter-efficient reasoning system,” in *ArXiv*, 2025.

* Equal Contribution

RESEARCH AND INDUSTRIAL EXPERIENCE

Empirical Analysis of SFT for Large Reasoning Models [6]

Jun 2025 - Sep 2025

Intern, Supervisor: Prof. Zhiting Hu

Institute of Foundation Models – MBZUAI / LLM 360

- My role in this project includes data pre-processing, model training, and technical report writing.
- Engineered and built the large-scale supervised fine-tuning (SFT) pipeline for models up to 70B parameters (e.g., LLaMA-3.1-70B, Qwen2.5-32B) on multiple GPU cluster.
- Achieved good results on challenging reasoning benchmarks by leveraging this pipeline, scoring 57.3 on LiveCodeBench (code generation) and 72.1 on AIME 2025 (math reasoning).
- Investigated key issues in SFT, including training loss behavior, data selection and mixing strategies, and the relationship between training dynamics and model performance; documented insights and best practices for future reference.

Memory LLM and Agents Benchmarking and Construction [3]–[5] Oct 2024 - Jun 2025
CSE Research Course, Supervisors: Yu Wang and Prof. Julian McAuley UC San Diego

- Led the development of **MemoryAgentBench**, a comprehensive benchmark designed to systematically assess the long-term memory of LLM agents via multi-turn interactions, with evaluation criteria based on principles of cognitive science.
- Designed and implemented the evaluation framework for M+, and conducted systematic benchmarking against baseline methods.
- Contributed to training Memory Alpha in long dialogue history settings, achieving 52% test accuracy on LongMemEval(S*) using only one-third of the full context window.
- Co-authored three research papers: MemoryAgentBench, RL Memory Agents and the M+ (ICML 2025).

De-biased Model Diagnosis Based on Eigspectrum and MP-Law [1] Nov 2024 - Apr 2025
Supervisors: Prof. Yaoqing Yang Dartmouth College

- Theorized and empirically validated that existing eigspectrum analysis is biased by weight matrix **aspect ratios**, a phenomenon explained by the **Marchenko-Pastur (MP) law**, leading to inaccurate layer diagnostics.
- Developed **FARMS**, a novel subsampling method grounded in **Random Matrix Theory (RMT)**, to normalize weight matrices to a fixed aspect ratio, enabling an unbiased, size-invariant evaluation of layer training quality.
- Validated the method’s effectiveness across diverse domains (**LLM Pruning, CV, SciML**), consistently **outperforming state-of-the-art** layer-wise optimization methods.
- First-authored a research paper accepted to the **ICML 2025**.

Layer-wise Optimization on Low-data SFT and Training [2] Jun 2023 - Jun 2024
Supervisors: Dr. Pu Ren and Prof. Yaoqing Yang Dartmouth College

- Utilized the **spectral analysis of model weights**, based on Heavy-Tailed Self-Regularization (HT-SR) theory, to establish a quantitative link between data scarcity and imbalanced layer-wise training, providing a theoretical foundation for designing superior model diagnosis metrics and optimization methods.
- Developed a dynamic **layer-wise learning rate** scheduling algorithm to rebalance training quality across layers, overcoming limitations of prior optimizers in NLP training scenarios.
- Validated the method’s effectiveness through extensive experiments on diverse models and benchmarks, improving LLM test accuracy **by 2-10%** in low-data SFT and leading to a **co-first** authored publication at EMNLP 2024 (**Oral Presentation**).
- Investigated the framework’s generalizability by extending the layer-wise method to tune **other hyperparameters** (e.g., weight decay, dropout), and validated its effectiveness on image classification with ResNet/ViT models on the CIFAR-100 dataset.
- Involved in the rebuttal process and supplementing experimental data for the paper accepted by NeurIPS 2023.

COMPETITION AWARDS / SERVICE

Second Prize at 17th National College Students Intelligent Car Race (Top-7% among 2771 teams), awarded by Chinese Automation Association, Summer 2022, Project Leader.
Reviewer for ICLR 2026, ARR 2025 (July / Oct.), and Workshops.
Volunteer for ICML 2025

TECHNICAL SKILLS

Programming Languages
Machine Learning Tools

Python, C/C++, SQL, Bash, Matlab
PyTorch, Hugging Face Transformers, LLaMA Factory, Verl